

# Project\_code

December 2, 2025

```
[60]: import numpy as np
import pandas as pd
import zipfile

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import (
    accuracy_score,
    precision_recall_fscore_support,
    roc_auc_score,
    confusion_matrix
)
from sklearn.dummy import DummyClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.cluster import KMeans
from sklearn.impute import SimpleImputer
```

```
[61]: import matplotlib.pyplot as plt

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from tensorflow.keras.optimizers import Adam
```

```
[62]: ZIP_PATH = "archive (16).zip"           # your uploaded zip file
CSV_NAME = "flight_data_2018_to_2022.csv"   # inside the zip

with zipfile.ZipFile(ZIP_PATH, "r") as z:
    with z.open(CSV_NAME) as f:
        df = pd.read_csv(f, low_memory=False)
```

```
print("Full shape:", df.shape)
df.head()
```

Full shape: (563737, 120)

```
[62]:    Year   Quarter   Month   DayofMonth   DayOfWeek   FlightDate   \
0   2022        1        1          6            4  2022-01-06
1   2022        1        1          6            4  2022-01-06
2   2022        1        1          6            4  2022-01-06
3   2022        1        1          6            4  2022-01-06
4   2022        1        1          6            4  2022-01-06

      Marketing_Airline_Network Operated_or_Branded_Code_Share_Partners   \
0                           DL                               DL
1                           DL                               DL
2                           DL                               DL
3                           DL                               DL
4                           DL                               DL

      DOT_ID_Marketing_Airline IATA_Code_Marketing_Airline   ...  Div5Airport   \
0             19790                      DL   ...       NaN
1             19790                      DL   ...       NaN
2             19790                      DL   ...       NaN
3             19790                      DL   ...       NaN
4             19790                      DL   ...       NaN

      Div5AirportID  Div5AirportSeqID  Div5WheelsOn  Div5TotalGTime   \
0           NaN           NaN           NaN           NaN
1           NaN           NaN           NaN           NaN
2           NaN           NaN           NaN           NaN
3           NaN           NaN           NaN           NaN
4           NaN           NaN           NaN           NaN

      Div5LongestGTime  Div5WheelsOff  Div5TailNum Duplicate  Unnamed: 119
0           NaN           NaN           NaN           N       NaN
1           NaN           NaN           NaN           N       NaN
2           NaN           NaN           NaN           N       NaN
3           NaN           NaN           NaN           N       NaN
4           NaN           NaN           NaN           N       NaN

[5 rows x 120 columns]
```

```
[63]: # Keep only rows where ArrDel15 is not missing
df = df[df["ArrDel15"].notna()].copy()

# Make sure ArrDel15 is int 0/1
df["ArrDel15"] = df["ArrDel15"].astype(int)
```

```

print("Shape after dropping NaN ArrDel15:", df.shape)
print("Delay rate (ArrDel15=1):", df["ArrDel15"].mean())

```

Shape after dropping NaN ArrDel15: (526894, 120)  
Delay rate (ArrDel15=1): 0.19456475116437083

```

[64]: def hhmm_to_minutes(x):
    """
    Convert HHMM numeric time (e.g., 1126) to minutes since midnight.
    Returns NaN if invalid.
    """
    try:
        x_int = int(x)
    except (ValueError, TypeError):
        return np.nan
    hour = x_int // 100
    minute = x_int % 100
    if hour < 0 or hour > 23 or minute < 0 or minute > 59:
        return np.nan
    return hour * 60 + minute

df["CRSDepMinutes"] = df["CRSDepTime"].apply(hhmm_to_minutes)

# Features (only pre-departure info)
feature_cols = [
    "Month",
    "DayOfWeek",
    "CRSDepMinutes",
    "Distance",
    "Origin",
    "Dest",
    "IATA_Code_Marketing_Airline"
]
target_col = "ArrDel15"

data = df[feature_cols + [target_col]].copy()
print("Selected columns:", data.columns.tolist())
data.head()

```

Selected columns: ['Month', 'DayOfWeek', 'CRSDepMinutes', 'Distance', 'Origin', 'Dest', 'IATA\_Code\_Marketing\_Airline', 'ArrDel15']

	Month	DayOfWeek	CRSDepMinutes	Distance	Origin	Dest	\
1	1	4	991	581.0	ATL	FLL	
2	1	4	1171	581.0	FLL	ATL	
3	1	4	624	680.0	FLL	RDU	

4	1	4	677	341.0	ATL	JAN
5	1	4	757	341.0	JAN	ATL
		IATA_Code_Marketing_Airline	ArrDel15			
1			DL	0		
2			DL	0		
3			DL	0		
4			DL	0		
5			DL	0		

## 0.1 overall delay rate from data (EDA)

```
[65]: # EDA Cell 1: Overall delay rate in the sampled data
overall_delay_rate = data["ArrDel15"].mean()
print(f"Overall delay rate (ArrDel15=1): {overall_delay_rate:.3f}")
```

Overall delay rate (ArrDel15=1): 0.195

## 0.2 delay rate by airline + bar plot

```
[66]: # EDA Cell 2: Delay rate by airline (IATA_Code_Marketing_Airline)

delay_by_airline = (
    data.groupby("IATA_Code_Marketing_Airline")["ArrDel15"]
        .mean()
        .sort_values(ascending=False)
)

print("Top 10 airlines by delay rate:")
print(delay_by_airline.head(10))


# Plot 1: Bar chart for top 10 airlines by delay rate
plt.figure(figsize=(8, 4))
delay_by_airline.head(10).plot(kind="bar")
plt.xlabel("Airline code")
plt.ylabel("Average delay rate (ArrDel15)")
plt.title("Delay rate by airline (top 10)")
plt.tight_layout()
plt.show()
```

Top 10 airlines by delay rate:  
IATA\_Code\_Marketing\_Airline  
B6 0.320185  
G4 0.284138  
F9 0.272767

```

NK    0.232032
AS    0.226563
UA    0.217331
HA    0.189539
WN    0.184186
AA    0.170686
DL    0.164472
Name: ArrDel15, dtype: float64

```



### 0.3 delay rate by month + bar plot


```
[67]: # EDA Cell 3: Delay rate by month (seasonality)

delay_by_month = data.groupby("Month")["ArrDel15"].mean()
print("Delay rate by month:")
print(delay_by_month)

# Optional small plot for your understanding (you can screenshot if you like)
plt.figure(figsize=(6, 4))
delay_by_month.plot(kind="bar")
plt.xlabel("Month")
plt.ylabel("Average delay rate (ArrDel15)")
plt.title("Delay rate by month")
plt.tight_layout()
plt.show()
```

Delay rate by month:  
Month

```
1      0.194565
Name: ArrDel15, dtype: float64
```



```
[68]: USE_SAMPLE = True          # set False to use all data
N_SAMPLES = 150_000            # adjust based on RAM / speed

if USE_SAMPLE and len(data) > N_SAMPLES:
    data = data.sample(n=N_SAMPLES, random_state=42).copy()
    print("Using sample:", data.shape)
else:
    print("Using full dataset:", data.shape)
```

```
Using sample: (150000, 8)
```

#### 0.4 define X, y and train\_test\_split (80/20, stratified)

```
[69]: X = data[feature_cols]
y = data[target_col]

X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,
    stratify=y,
```

```

    random_state=42
)

print("Train size:", X_train.shape, " Test size:", X_test.shape)
print("Train delay rate:", y_train.mean())
print("Test delay rate :", y_test.mean())

```

Train size: (120000, 7) Test size: (30000, 7)  
 Train delay rate: 0.1930333333333333  
 Test delay rate : 0.1930333333333333

## 0.5 Preprocessing

```
[70]: numeric_features = ["Month", "DayOfWeek", "CRSDepMinutes", "Distance"]
categorical_features = ["Origin", "Dest", "IATA_Code_Marketing_Airline"]

def make_preprocess():
    numeric_transformer = Pipeline(steps=[
        ("imputer", SimpleImputer(strategy="median")),
        ("scaler", StandardScaler())
    ])

    categorical_transformer = Pipeline(steps=[
        ("imputer", SimpleImputer(strategy="most_frequent")),
        ("onehot", OneHotEncoder(handle_unknown="ignore"))
    ])

    preprocessor = ColumnTransformer(
        transformers=[
            ("num", numeric_transformer, numeric_features),
            ("cat", categorical_transformer, categorical_features)
        ]
    )
    return preprocessor
```

```
[71]: def evaluate_model(name, y_true, y_pred, y_proba=None):
    acc = accuracy_score(y_true, y_pred)
    precision, recall, f1, _ = precision_recall_fscore_support(
        y_true, y_pred,
        average="binary",
        pos_label=1,
        zero_division=0
    )
    if y_proba is not None:
        try:
            auc = roc_auc_score(y_true, y_proba)
```

```

        except ValueError:
            auc = np.nan
    else:
        auc = np.nan

    print(f"\n==== {name} ===")
    print(f"Accuracy : {acc:.4f}")
    print(f"Precision: {precision:.4f}")
    print(f"Recall   : {recall:.4f}")
    print(f"F1-score : {f1:.4f}")
    print(f"ROC-AUC  : {auc:.4f}")
    print("Confusion matrix:")
    print(confusion_matrix(y_true, y_pred))

    return {
        "model": name,
        "accuracy": acc,
        "precision": precision,
        "recall": recall,
        "f1": f1,
        "roc_auc": auc
    }

results = []

```

## 0.6 Baseline models

```
[72]: dummy_clf = DummyClassifier(strategy="most_frequent")
dummy_clf.fit(X_train, y_train)
y_pred_dummy = dummy_clf.predict(X_test)

results.append(evaluate_model("Majority Class (Dummy)", y_test, y_pred_dummy))
```

```

==== Majority Class (Dummy) ===
Accuracy : 0.8070
Precision: 0.0000
Recall   : 0.0000
F1-score : 0.0000
ROC-AUC  : nan
Confusion matrix:
[[24209      0]
 [ 5791      0]]
```

```
[73]: # Logistic Regression
log_reg = Pipeline(steps=[
```

```

        ("preprocess", make_preprocess()),
        ("clf", LogisticRegression(
            max_iter=1000,
            class_weight="balanced",
            n_jobs=-1
        )))
    ])

log_reg.fit(X_train, y_train)
y_pred_lr = log_reg.predict(X_test)
y_proba_lr = log_reg.predict_proba(X_test)[:, 1]
results.append(evaluate_model("Logistic Regression", y_test, y_pred_lr, y_proba_lr))

# Decision Tree
dt_clf = Pipeline(steps=[
    ("preprocess", make_preprocess()),
    ("clf", DecisionTreeClassifier(
        max_depth=10,
        class_weight="balanced",
        random_state=42
    )))
])

dt_clf.fit(X_train, y_train)
y_pred_dt = dt_clf.predict(X_test)
y_proba_dt = dt_clf.predict_proba(X_test)[:, 1]
results.append(evaluate_model("Decision Tree", y_test, y_pred_dt, y_proba_dt))

```

==== Logistic Regression ===

Accuracy : 0.5977

Precision: 0.2592

Recall : 0.5833

F1-score : 0.3589

ROC-AUC : 0.6277

Confusion matrix:

[[14553 9656]	
[ 2413 3378]]	

==== Decision Tree ===

Accuracy : 0.5751

Precision: 0.2535

Recall : 0.6177

F1-score : 0.3595

ROC-AUC : 0.6232

Confusion matrix:

```
[[13676 10533]
 [ 2214  3577]]
```

```
[74]: # Random Forest
rf_clf = Pipeline(steps=[
    ("preprocess", make_preprocess()),
    ("clf", RandomForestClassifier(
        n_estimators=200,
        max_depth=15,
        class_weight="balanced_subsample",
        n_jobs=-1,
        random_state=42
    ))
])

rf_clf.fit(X_train, y_train)
y_pred_rf = rf_clf.predict(X_test)
y_proba_rf = rf_clf.predict_proba(X_test)[:, 1]
results.append(evaluate_model("Random Forest", y_test, y_pred_rf, y_proba_rf))
```

==== Random Forest ===

Accuracy : 0.6366

Precision: 0.2758

Recall : 0.5431

F1-score : 0.3658

ROC-AUC : 0.6437

Confusion matrix:

```
[[15952  8257]
```

```
 [ 2646  3145]]
```

```
[75]: from sklearn.svm import LinearSVC

# SVM using LinearSVC (faster, no probability calibration needed)
svm_clf = Pipeline(steps=[
    ("preprocess", make_preprocess()),
    ("clf", LinearSVC(
        class_weight="balanced",
        random_state=42
    ))
])

# Train
svm_clf.fit(X_train, y_train)

# Class predictions (0/1)
y_pred_svm = svm_clf.predict(X_test)
```

```

# Use decision_function scores as "probabilities" for ROC-AUC
# (roc_auc_score works with any continuous scores, not only true probabilities)
y_score_svm = svm_clf.decision_function(X_test)

# Evaluate - pass y_score_sum in place of y_proba
results.append(evaluate_model(
    "SVM (LinearSVC, decision_function)",
    y_test,
    y_pred_svm,
    y_score_svm
))

```

```

==== SVM (LinearSVC, decision_function) ====
Accuracy : 0.5976
Precision: 0.2594
Recall   : 0.5847
F1-score : 0.3594
ROC-AUC  : 0.6270
Confusion matrix:
[[14543  9666]
 [ 2405  3386]]

```

## 0.7 Clustering + CSE

```
[76]: cluster_numeric_features = ["Month", "DayOfWeek", "CRSDepMinutes", "Distance"]

X_train_num = X_train[cluster_numeric_features].copy()
X_test_num = X_test[cluster_numeric_features].copy()

cluster_scaler = StandardScaler()
X_train_num_scaled = cluster_scaler.fit_transform(X_train_num)
X_test_num_scaled = cluster_scaler.transform(X_test_num)
```

```
[77]: K = 4 # you can try 3, 4, 5 and compare

kmeans = KMeans(n_clusters=K, random_state=42, n_init=10)
train_clusters = kmeans.fit_predict(X_train_num_scaled)
test_clusters = kmeans.predict(X_test_num_scaled)

unique, counts = np.unique(train_clusters, return_counts=True)
print("Train cluster sizes:")
for u, c in zip(unique, counts):
    print(f"Cluster {u}: {c} samples")
```

```
Train cluster sizes:  
Cluster 0: 32743 samples  
Cluster 1: 33530 samples  
Cluster 2: 13599 samples  
Cluster 3: 40128 samples
```

```
[78]: cluster_models = {}  
MIN_CLUSTER_SIZE = 1000    # skip tiny clusters if you want  
  
for c in range(K):  
    idx = np.where(train_clusters == c)[0]  
    if len(idx) < MIN_CLUSTER_SIZE:  
        print(f"Skipping cluster {c} (too small: {len(idx)} samples)")  
        continue  
  
    X_train_c = X_train.iloc[idx]  
    y_train_c = y_train.iloc[idx]  
  
    print(f"Training model for cluster {c} with {len(idx)} samples")  
  
    model_c = Pipeline(steps=[  
        ("preprocess", make_preprocess()),  
        ("clf", RandomForestClassifier(  
            n_estimators=200,  
            max_depth=15,  
            class_weight="balanced_subsample",  
            n_jobs=-1,  
            random_state=42  
        ))  
    ])  
  
    model_c.fit(X_train_c, y_train_c)  
    cluster_models[c] = model_c  
  
print("Trained cluster models:", list(cluster_models.keys()))
```

```
Training model for cluster 0 with 32743 samples  
Training model for cluster 1 with 33530 samples  
Training model for cluster 2 with 13599 samples  
Training model for cluster 3 with 40128 samples  
Trained cluster models: [0, 1, 2, 3]
```

## 0.8 Confusion matrix plot for global Random Forest


```
[79]: from sklearn.metrics import ConfusionMatrixDisplay

cm_rf = confusion_matrix(y_test, y_pred_rf)

disp_rf = ConfusionMatrixDisplay(
    confusion_matrix=cm_rf,
    display_labels=["On time (0)", "Delayed (1)"]
)

fig, ax = plt.subplots(figsize=(4, 4))
disp_rf.plot(ax=ax)

plt.tight_layout()
plt.show()
```



## 0.9 Confusion matrix plot for CSE

[80]: # 0.10 Cluster-Specific Ensemble (CSE) + Confusion Matrix Plot 2

```
# Start from global RF predictions
y_pred_cse = rf_clf.predict(X_test)
y_proba_cse = rf_clf.predict_proba(X_test)[:, 1]


# Override with cluster-specific models
for c, model_c in cluster_models.items():
    idx = np.where(test_clusters == c)[0]
    if len(idx) == 0:
        continue
    X_test_c = X_test.iloc[idx]
    preds_c = model_c.predict(X_test_c)
    proba_c = model_c.predict_proba(X_test_c)[:, 1]

    y_pred_cse[idx] = preds_c
    y_proba_cse[idx] = proba_c

# Now plot confusion matrix
cm_cse = confusion_matrix(y_test, y_pred_cse)

disp_cse = ConfusionMatrixDisplay(
    confusion_matrix=cm_cse,
    display_labels=["On time (0)", "Delayed (1)"]
)

fig, ax = plt.subplots(figsize=(4, 4))
disp_cse.plot(ax=ax)
plt.tight_layout()
plt.show()
```



## 0.10 Model Performance Comparison

```
[81]: # Start from global Random Forest predictions (fallback)
y_pred_cse = rf_clf.predict(X_test)
y_proba_cse = rf_clf.predict_proba(X_test)[:, 1]

# Override with cluster-specific models when available
for c, model_c in cluster_models.items():
    idx = np.where(test_clusters == c)[0]
    if len(idx) == 0:
        continue

    X_test_c = X_test.iloc[idx]
    preds_c = model_c.predict(X_test_c)
    proba_c = model_c.predict_proba(X_test_c)[:, 1]

    y_pred_cse[idx] = preds_c
    y_proba_cse[idx] = proba_c

# Evaluate CSE model
results.append(evaluate_model("Cluster-Specific Ensemble (CSE)", y_test,
                             y_pred_cse, y_proba_cse))
results_df = pd.DataFrame(results).set_index("model")
```

```
results_df
```

```
==== Cluster-Specific Ensemble (CSE) ====
Accuracy : 0.6534
Precision: 0.2783
Recall   : 0.4992
F1-score  : 0.3574
ROC-AUC   : 0.6375
Confusion matrix:
[[16711  7498]
 [ 2900  2891]]
```

```
[81]:
```

model	accuracy	precision	recall	f1	\
Majority Class (Dummy)	0.806967	0.000000	0.000000	0.000000	
Logistic Regression	0.597700	0.259168	0.583319	0.358884	
Decision Tree	0.575100	0.253508	0.617683	0.359479	
Random Forest	0.636567	0.275829	0.543084	0.365847	
SVM (LinearSVC, decision_function)	0.597633	0.259424	0.584700	0.359391	
Cluster-Specific Ensemble (CSE)	0.653400	0.278275	0.499223	0.357355	

model	roc_auc
Majority Class (Dummy)	NaN
Logistic Regression	0.627707
Decision Tree	0.623248
Random Forest	0.643703
SVM (LinearSVC, decision_function)	0.626968
Cluster-Specific Ensemble (CSE)	0.637490

## 0.11 Preprocess for neural net, build & train LSTM

```
[82]: # Use the same preprocessing as other models, but get a dense array
preprocess_nn = make_preprocess() # reuses your numeric + categorical pipeline

X_train_nn = preprocess_nn.fit_transform(X_train)
X_test_nn = preprocess_nn.transform(X_test)

# If the result is sparse (because of one-hot), convert to dense
if hasattr(X_train_nn, "toarray"):
    X_train_nn = X_train_nn.toarray()
    X_test_nn = X_test_nn.toarray()

print("Neural net feature shape:", X_train_nn.shape)
```

```

# Reshape for LSTM: (samples, timesteps, features)
# We treat each flight as a sequence of length 1 with many features
n_features = X_train_nn.shape[1]
X_train_lstm = X_train_nn.reshape(-1, 1, n_features)
X_test_lstm = X_test_nn.reshape(-1, 1, n_features)

# Define a simple LSTM model
lstm_model = Sequential([
    LSTM(32, input_shape=(1, n_features)),
    Dense(16, activation="relu"),
    Dense(1, activation="sigmoid")
])
lstm_model.compile(
    loss="binary_crossentropy",
    optimizer=Adam(learning_rate=1e-3),
    metrics=["accuracy"]
)
lstm_model.summary()

# Train the LSTM - you can increase epochs later if it's fast
history_lstm = lstm_model.fit(
    X_train_lstm, y_train,
    validation_split=0.2,
    epochs=5,
    batch_size=256,
    verbose=1
)

# Predict on the test set
y_proba_lstm = lstm_model.predict(X_test_lstm).ravel()
y_pred_lstm = (y_proba_lstm >= 0.5).astype(int)

```

Neural net feature shape: (120000, 744)

```

/usr/local/lib/python3.11/site-packages/keras/src/layers/rnn/rnn.py:204:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.

```

```
super().__init__(**kwargs)
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
lstm_3 (LSTM)	(None, 32)	99,456

dense_6 (Dense)	(None, 16)	528
dense_7 (Dense)	(None, 1)	17

Total params: 100,001 (390.63 KB)

Trainable params: 100,001 (390.63 KB)

Non-trainable params: 0 (0.00 B)

```
Epoch 1/5
375/375      3s 4ms/step -
accuracy: 0.7972 - loss: 0.5441 - val_accuracy: 0.8089 - val_loss: 0.4709
Epoch 2/5
375/375      1s 3ms/step -
accuracy: 0.8053 - loss: 0.4745 - val_accuracy: 0.8090 - val_loss: 0.4686
Epoch 3/5
375/375      1s 3ms/step -
accuracy: 0.8067 - loss: 0.4710 - val_accuracy: 0.8092 - val_loss: 0.4671
Epoch 4/5
375/375      1s 3ms/step -
accuracy: 0.8079 - loss: 0.4657 - val_accuracy: 0.8096 - val_loss: 0.4668
Epoch 5/5
375/375      1s 3ms/step -
accuracy: 0.8063 - loss: 0.4666 - val_accuracy: 0.8100 - val_loss: 0.4666
938/938      1s 1ms/step
```

## 0.12 LSTM to our results table

```
[83]: # Add LSTM to the results list and rebuild results_df
results.append(
    evaluate_model("LSTM", y_test, y_pred_lstm, y_proba_lstm)
)

results_df = pd.DataFrame(results).set_index("model")
results_df
```

```
==== LSTM ====
Accuracy : 0.8077
Precision: 0.6146
Recall   : 0.0102
F1-score : 0.0200
```

```
ROC-AUC : 0.6427
```

```
Confusion matrix:
```

```
[[24172 37]
 [ 5732 59]]
```

[83]: model	accuracy	precision	recall	f1 \
Majority Class (Dummy)	0.806967	0.000000	0.000000	0.000000
Logistic Regression	0.597700	0.259168	0.583319	0.358884
Decision Tree	0.575100	0.253508	0.617683	0.359479
Random Forest	0.636567	0.275829	0.543084	0.365847
SVM (LinearSVC, decision_function)	0.597633	0.259424	0.584700	0.359391
Cluster-Specific Ensemble (CSE)	0.653400	0.278275	0.499223	0.357355
LSTM	0.807700	0.614583	0.010188	0.020044

[83]: model	roc_auc
Majority Class (Dummy)	NaN
Logistic Regression	0.627707
Decision Tree	0.623248
Random Forest	0.643703
SVM (LinearSVC, decision_function)	0.626968
Cluster-Specific Ensemble (CSE)	0.637490
LSTM	0.642652

```
[ ]:
```